

```
In [ ]: import matplotlib.pyplot as plt
import pandas as pd
from sklearn.decomposition import PCA
from sklearn.preprocessing import StandardScaler
from sklearn.model_selection import train_test_split
from sklearn.ensemble import RandomForestClassifier
from sklearn.svm import SVC
from sklearn.linear_model import LogisticRegression
from sklearn.metrics import confusion_matrix, classification_report
from sklearn.model_selection import GridSearchCV
from sklearn.neural_network import MLPClassifier
from sklearn.feature_selection import SelectKBest, chi2
from collections import Counter
import seaborn as sns
from imblearn.over_sampling import RandomOverSampler, SMOTE
from imblearn.under_sampling import RandomUnderSampler
import warnings
warnings.filterwarnings("ignore")
```

1. Data Preprocessing

Importing dataset & Checking for missing data

```
In [ ]: headerList = ['fixed acidity', 'volatile acidity', 'citric acid', 'residual sugar', 'chlorides', 'free sulfur dioxide', 'total sulfur dioxide', 'density', 'pH', 'sulphates',
                        'alcohol']

wineData = pd.read_csv('winequality-red.csv', header = 0, names = headerList, sep=";")
print(wineData.head(10))
print(wineData)

#Summarative functions
wineData.dtypes
wineData.describe()
wineData.info()

#Plotting histogram of each variable
wineData.hist(alpha=0.5, figsize=(15, 10))
plt.tight_layout()
plt.show()

for h in headerList:
    wineData[h] = pd.to_numeric(wineData[h], errors='coerce')

print("\nChecking for null values: \n")
wineData.isna().sum()
wineData = wineData.fillna(0)
print("\nChecking for null values after using fillna(): \n")
wineData.isna().sum()
```

	fixed acidity	volatile acidity	citric acid	residual sugar	chlorides	
0	7.4	0.70	0.00	1.9	0.076	\
1	7.8	0.88	0.00	2.6	0.098	
2	7.8	0.76	0.04	2.3	0.092	
3	11.2	0.28	0.56	1.9	0.075	
4	7.4	0.70	0.00	1.9	0.076	
5	7.4	0.66	0.00	1.8	0.075	
6	7.9	0.60	0.06	1.6	0.069	
7	7.3	0.65	0.00	1.2	0.065	
8	7.8	0.58	0.02	2.0	0.073	
9	7.5	0.50	0.36	6.1	0.071	

	free sulfur dioxide	total sulfur dioxide	density	pH	sulphates	
0	11.0	34.0	0.9978	3.51	0.56	\
1	25.0	67.0	0.9968	3.20	0.68	
2	15.0	54.0	0.9970	3.26	0.65	
3	17.0	60.0	0.9980	3.16	0.58	
4	11.0	34.0	0.9978	3.51	0.56	
5	13.0	40.0	0.9978	3.51	0.56	
6	15.0	59.0	0.9964	3.30	0.46	
7	15.0	21.0	0.9946	3.39	0.47	
8	9.0	18.0	0.9968	3.36	0.57	
9	17.0	102.0	0.9978	3.35	0.80	

	alcohol	quality
0	9.4	5
1	9.8	5
2	9.8	5
3	9.8	6
4	9.4	5
5	9.4	5
6	9.4	5
7	10.0	7
8	9.5	7
9	10.5	5

	fixed acidity	volatile acidity	citric acid	residual sugar	chlorides	
0	7.4	0.700	0.00	1.9	0.076	\
1	7.8	0.880	0.00	2.6	0.098	
2	7.8	0.760	0.04	2.3	0.092	
3	11.2	0.280	0.56	1.9	0.075	
4	7.4	0.700	0.00	1.9	0.076	
...
1594	6.2	0.600	0.08	2.0	0.090	
1595	5.9	0.550	0.10	2.2	0.062	
1596	6.3	0.510	0.13	2.3	0.076	
1597	5.9	0.645	0.12	2.0	0.075	
1598	6.0	0.310	0.47	3.6	0.067	

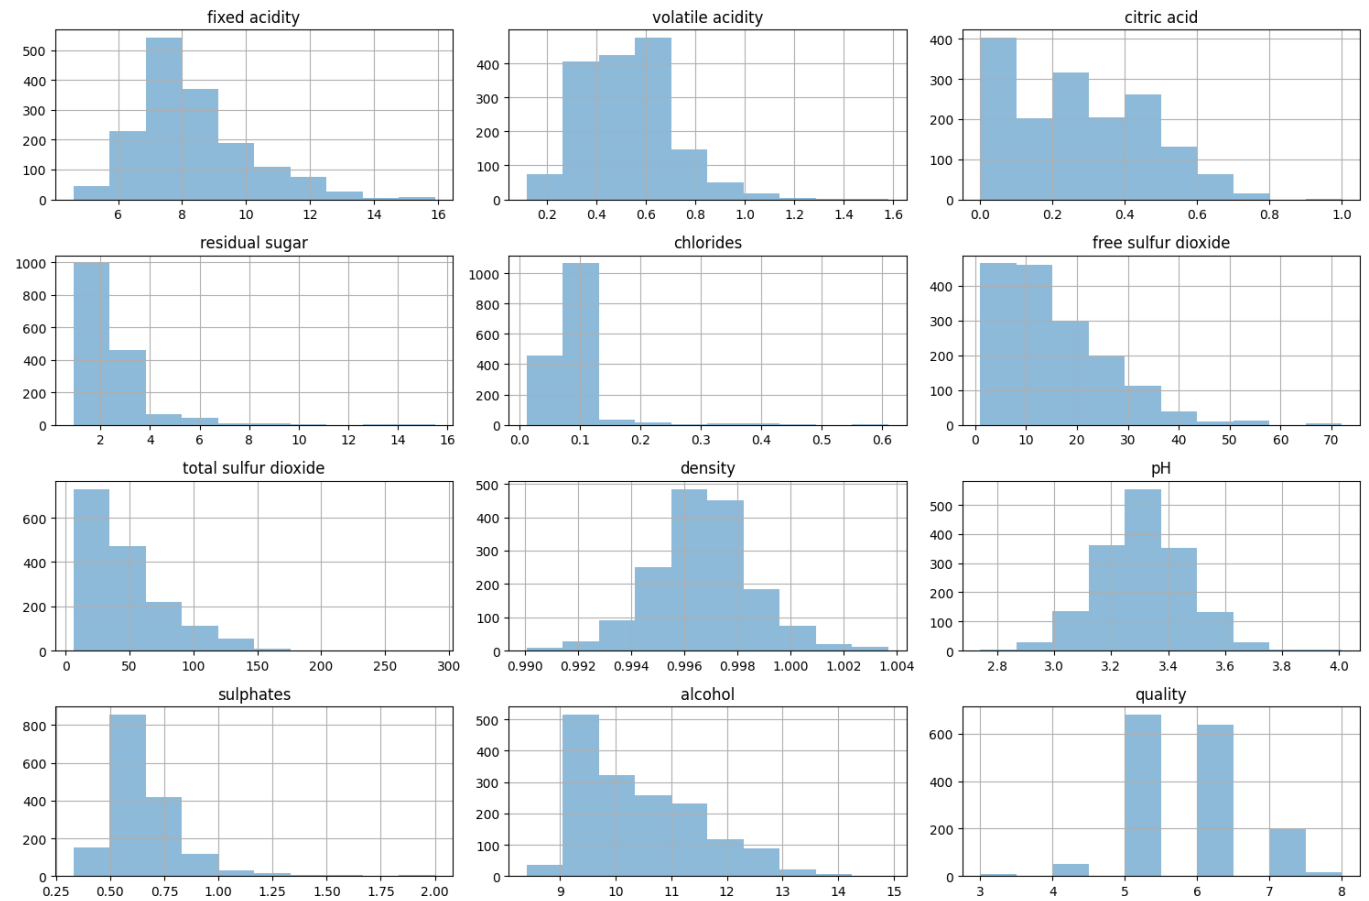
	free sulfur dioxide	total sulfur dioxide	density	pH	sulphates	
0	11.0	34.0	0.99780	3.51	0.56	\
1	25.0	67.0	0.99680	3.20	0.68	
2	15.0	54.0	0.99700	3.26	0.65	
3	17.0	60.0	0.99800	3.16	0.58	
4	11.0	34.0	0.99780	3.51	0.56	
...
1594	32.0	44.0	0.99490	3.45	0.58	
1595	39.0	51.0	0.99512	3.52	0.76	
1596	29.0	40.0	0.99574	3.42	0.75	
1597	32.0	44.0	0.99547	3.57	0.71	
1598	18.0	42.0	0.99549	3.39	0.66	

	alcohol	quality
0	9.4	5
1	9.8	5
2	9.8	5
3	9.8	6
4	9.4	5
...
1594	10.5	5
1595	11.2	6
1596	11.0	6
1597	10.2	5
1598	11.0	6

```
[1599 rows x 12 columns]
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 1599 entries, 0 to 1598
Data columns (total 12 columns):
```

#	Column	Non-Null Count	Dtype
0	fixed acidity	1599 non-null	float64
1	volatile acidity	1599 non-null	float64
2	citric acid	1599 non-null	float64
3	residual sugar	1599 non-null	float64
4	chlorides	1599 non-null	float64
5	free sulfur dioxide	1599 non-null	float64
6	total sulfur dioxide	1599 non-null	float64
7	density	1599 non-null	float64
8	pH	1599 non-null	float64
9	sulphates	1599 non-null	float64
10	alcohol	1599 non-null	float64
11	quality	1599 non-null	int64

```
dtypes: float64(11), int64(1)
memory usage: 150.0 KB
```



Checking for null values:

Checking for null values after using fillna():

```
Out[ ]:
fixed acidity      0
volatile acidity   0
citric acid        0
residual sugar     0
chlorides          0
free sulfur dioxide 0
total sulfur dioxide 0
density            0
pH                0
sulphates          0
alcohol            0
quality            0
dtype: int64
```

Our data is imbalanced as seen from the histogram. We will adapt multiple strategies to address the issue.

2. Exploratory Data Analysis

1. Principal Component Analysis (PCA)

```
In [ ]:
#PCA
df_pca = wineData.copy()
X_pca = df_pca.loc[:, 'fixed acidity':'alcohol']
y_pca = df_pca['quality']


X_pca.tail()
X_pca = StandardScaler().fit_transform(X_pca)

#Fit PCA
pca = PCA(n_components=2)
X_pca = pca.fit_transform(X_pca)

X_pca.shape

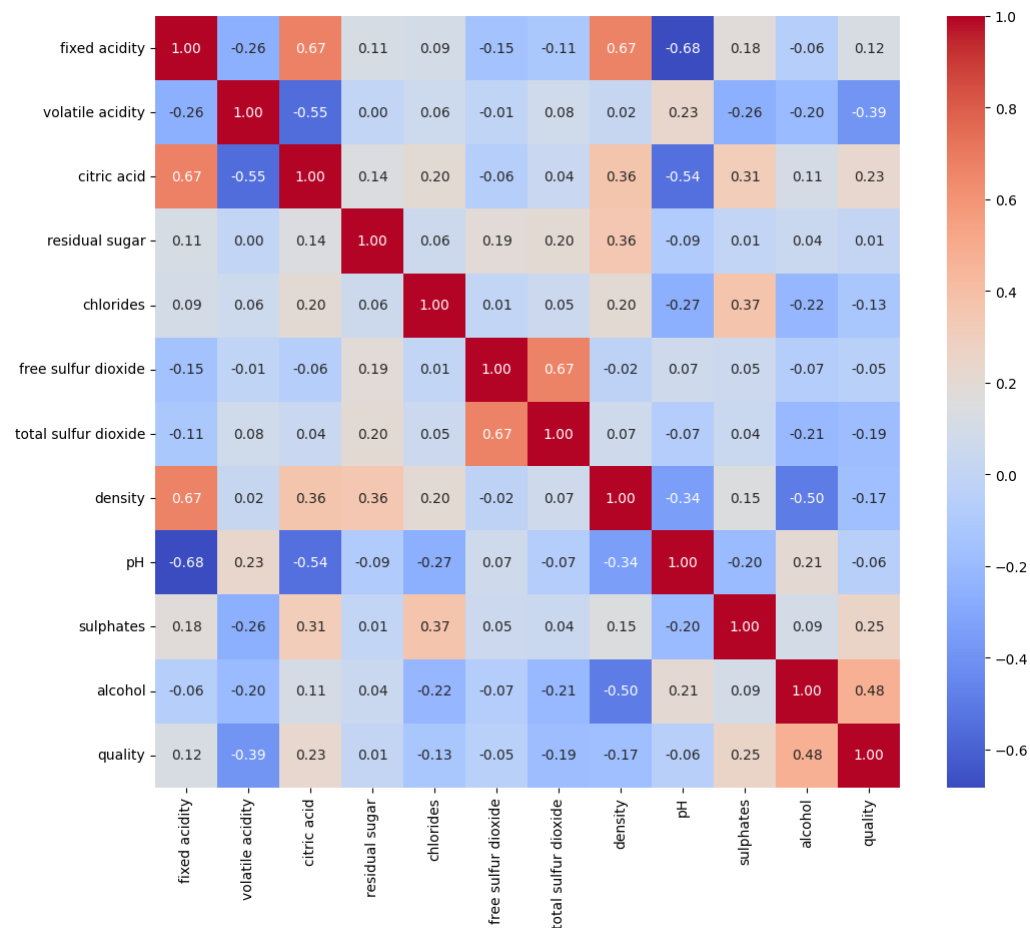
plt.figure(figsize=(8,6))
plt.scatter(X_pca[:,0],X_pca[:,1],c=y_pca,cmap='rainbow')
plt.xlabel('First principal component')
plt.ylabel('Second Principal Component')
plt.title("Using PCA to Visualize Classes")
plt.show()

print("components: ", pca.components_, "\n")
print("explained variance: ", pca.explained_variance_, "\n")
exp_var_rat = pca.explained_variance_ratio_
print("explained variance ratio: ", exp_var_rat)
```



2. Correlation Matrix

```
In [ ]: plt.figure(figsize=(12,10))
sns.heatmap(wineData.corr(),annot=True, cmap='coolwarm',fmt='.2f')
Out [ ]: <Axes: >
```



3. Univariate Selection

```
In [ ]: #Split data into training and test sets
X = wineData.loc[:, 'fixed acidity':'alcohol']
y = wineData['quality']

# apply SelectKBest class to extract best features
bestFeatures = SelectKBest(score_func=chi2, k=11)
bestFeaturesFit = bestFeatures.fit(X,y)
```

```
dfscores = pd.DataFrame(bestFeaturesFit.scores_)
dfcolumns = pd.DataFrame(X.columns)

# concatenate scores with predictor names
predScores = pd.concat([dfcolumns,dfscores],axis=1)
predScores.columns = ['Predictor','Score']
print(predScores.nlargest(11,'Score'))
```

	Predictor	Score
6	total sulfur dioxide	2755.557984
5	free sulfur dioxide	161.936036
10	alcohol	46.429892
1	volatile acidity	15.580289
2	citric acid	13.025665
0	fixed acidity	11.260652
9	sulphates	4.558488
3	residual sugar	4.123295
4	chlorides	0.752426
8	pH	0.154655
7	density	0.000230

Dropping features from univariate selection

We are dropping bottom features as they have very low predictor scores and to save computation

```
In [ ]: #Drop the bottom four features (smallest score)
wineData = wineData.drop(['density'], axis=1)
wineData = wineData.drop(['pH'], axis=1)
wineData = wineData.drop(['chlorides'], axis=1)
print(wineData)

X = wineData.loc[:, 'fixed acidity':'alcohol']
y = wineData['quality']
```

	fixed acidity	volatile acidity	citric acid	residual sugar	
0	7.4	0.700	0.00	1.9	\
1	7.8	0.880	0.00	2.6	
2	7.8	0.760	0.04	2.3	
3	11.2	0.280	0.56	1.9	
4	7.4	0.700	0.00	1.9	
...	
1594	6.2	0.600	0.08	2.0	
1595	5.9	0.550	0.10	2.2	
1596	6.3	0.510	0.13	2.3	
1597	5.9	0.645	0.12	2.0	
1598	6.0	0.310	0.47	3.6	

	free sulfur dioxide	total sulfur dioxide	sulphates	alcohol	quality
0	11.0	34.0	0.56	9.4	5
1	25.0	67.0	0.68	9.8	5
2	15.0	54.0	0.65	9.8	5
3	17.0	60.0	0.58	9.8	6
4	11.0	34.0	0.56	9.4	5
...
1594	32.0	44.0	0.58	10.5	5
1595	39.0	51.0	0.76	11.2	6
1596	29.0	40.0	0.75	11.0	6
1597	32.0	44.0	0.71	10.2	5
1598	18.0	42.0	0.66	11.0	6

[1599 rows x 9 columns]

Addressing Imbalance in Class

First Strategy: Oversampling minority class

```
In [ ]: oversample = RandomOverSampler(sampling_strategy='minority')
X_over, y_over = oversample.fit_resample(X, y)
print("Before RandomOverSampler : ", Counter(y))
print("After RandomOverSampler : ", Counter(y_over))

Before RandomOverSampler : Counter({5: 681, 6: 638, 7: 199, 4: 53, 8: 18, 3: 10})
After RandomOverSampler : Counter({5: 681, 3: 681, 6: 638, 7: 199, 4: 53, 8: 18})
```

Second Strategy: Undersampling majority class

```
In [ ]: undersample = RandomUnderSampler(sampling_strategy='majority')
X_under, y_under = undersample.fit_resample(X, y)
print("Before RandomUnderSampler : ", Counter(y))
print("After RandomUnderSampler : ", Counter(y_under))

Before RandomUnderSampler : Counter({5: 681, 6: 638, 7: 199, 4: 53, 8: 18, 3: 10})
After RandomUnderSampler : Counter({6: 638, 7: 199, 4: 53, 8: 18, 3: 10, 5: 10})
```

Third Strategy: SMOTE


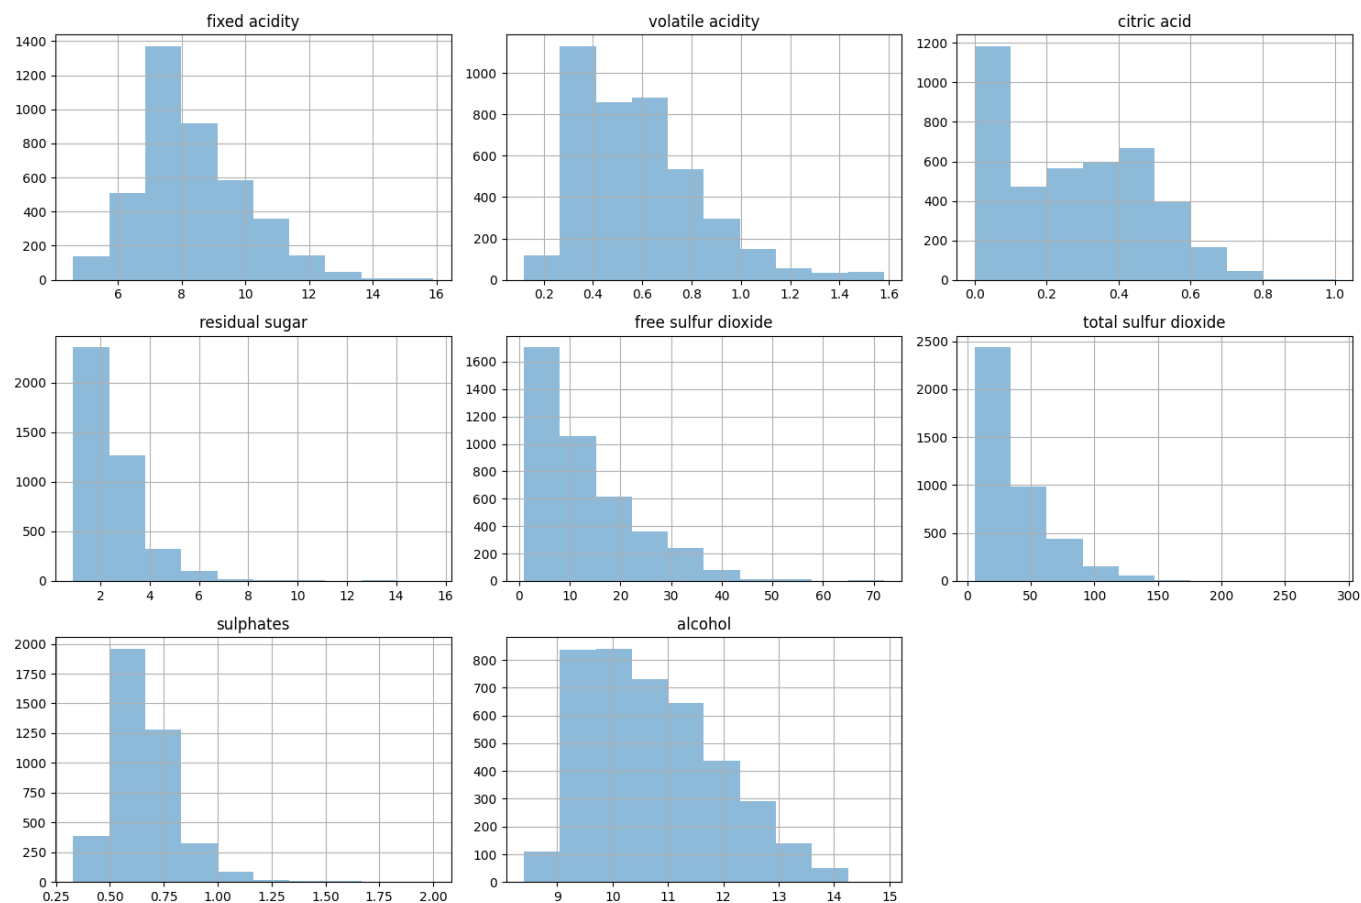
```
In [ ]: smoteOversample = SMOTE()
X_smote, y_smote = smoteOversample.fit_resample(X, y)

#Plotting histogram of each variable
X_smote.hist(alpha=0.5, figsize=(15, 10))

plt.tight_layout()
plt.show()

y_smote.hist(alpha=0.5, figsize=(15, 10))
plt.show()

from collections import Counter
print("Before SMOTE : ", Counter(y))
print("After SMOTE : ", Counter(y_smote))
```



Before SMOTE : Counter({5: 681, 6: 638, 7: 199, 4: 53, 8: 18, 3: 10})
 After SMOTE : Counter({5: 681, 6: 681, 7: 681, 4: 681, 8: 681, 3: 681})

Fourth Strategy: Data Imputation

Filling in data from missing classes - 0, 1, 2, 9, & 10 with fraud data

```
In [ ]: avgX = X.mean(axis=0)
dfImpute = pd.DataFrame(
    [[avgX[0], avgX[1], avgX[2], avgX[3], avgX[4], avgX[5], avgX[6], avgX[7], 0],
     [avgX[0], avgX[1], avgX[2], avgX[3], avgX[4], avgX[5], avgX[6], avgX[7], 1],
     [avgX[0], avgX[1], avgX[2], avgX[3], avgX[4], avgX[5], avgX[6], avgX[7], 2],
```

```

[avgX[0], avgX[1], avgX[2], avgX[3], avgX[4], avgX[5], avgX[6], avgX[7], 9],
[avgX[0], avgX[1], avgX[2], avgX[3], avgX[4], avgX[5], avgX[6], avgX[7], 10]],
columns=['fixed acidity', 'volatile acidity', 'citric acid', 'residual sugar', 'free sulfur dioxide', 'total sulfur dioxide', 'sulphates', 'alcohol',
])
dfImpute

```

```

Out [ ]:
   fixed acidity  volatile acidity  citric acid  residual sugar  free sulfur dioxide  total sulfur dioxide  sulphates  alcohol  quality
0      8.319637      0.527821      0.270976      2.538806      15.874922      46.467792      0.658149      10.422983      0
1      8.319637      0.527821      0.270976      2.538806      15.874922      46.467792      0.658149      10.422983      1
2      8.319637      0.527821      0.270976      2.538806      15.874922      46.467792      0.658149      10.422983      2
3      8.319637      0.527821      0.270976      2.538806      15.874922      46.467792      0.658149      10.422983      9
4      8.319637      0.527821      0.270976      2.538806      15.874922      46.467792      0.658149      10.422983      10

```

3. Comparing Machine Learning Models / Obtaining Baseline Accuracy

Modeling - Final data preparations

```

In [ ]:
# 1. Regular Data
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size = .2, random_state=10) #split the data
X_train.shape, y_train.shape, X_test.shape, y_test.shape
scaledData = StandardScaler()
X_train = scaledData.fit_transform(X_train)
X_test = scaledData.transform(X_test)

# 2. Oversampled Data
X_over_train, X_over_test, y_over_train, y_over_test = train_test_split(X_over, y_over, test_size = .2, random_state=10) #split the data
X_over_train.shape, y_over_train.shape, X_over_test.shape, y_over_test.shape
X_over_train = scaledData.fit_transform(X_over_train)
X_over_test = scaledData.transform(X_over_test)

# 3. Undersampled Data
X_under_train, X_under_test, y_under_train, y_under_test = train_test_split(X_under, y_under, test_size = .2, random_state=10) #split the data
X_under_train.shape, y_under_train.shape, X_under_test.shape, y_under_test.shape
X_under_train = scaledData.fit_transform(X_under_train)
X_under_test = scaledData.transform(X_under_test)

#4. SMOTE Data
X_smote_train, X_smote_test, y_smote_train, y_smote_test = train_test_split(X_smote, y_smote, test_size = .2, random_state=10) #split the data
X_smote_train.shape, y_smote_train.shape, X_smote_test.shape, y_smote_test.shape
X_smote_train = scaledData.fit_transform(X_smote_train)
X_smote_test = scaledData.transform(X_smote_test)

# 5. Imputed Data
X_impute = wineData.loc[:, 'fixed acidity':'alcohol']
y_impute = wineData['quality']
X_add = dfImpute.loc[:, 'fixed acidity':'alcohol']
y_add = dfImpute['quality']
X_impute_train, X_impute_test, y_impute_train, y_impute_test = train_test_split(X_impute, y_impute, test_size = .2, random_state=10) #split the data
X_impute_train = pd.concat([X_impute_train, X_add])
y_impute_train = pd.concat([y_impute_train, y_add])
X_impute_test = pd.concat([X_impute_test, X_add])
y_impute_test = pd.concat([y_impute_test, y_add])
X_impute_train.shape, y_impute_train.shape, X_impute_test.shape, y_impute_test.shape
X_impute_train = scaledData.fit_transform(X_impute_train)

def clas_report(X_train, y_train, X_test, y_test, model, title):
    model.fit(X_train, y_train)
    y_pred = model.predict(X_test)
    cm = confusion_matrix(y_test, y_pred)
    df_cm = pd.DataFrame(cm)
    sns.set(font_scale=1.2) # for label size
    sns.heatmap(df_cm, annot=True, annot_kws={"size": 16}) # font size
    plt.show()
    clas = classification_report(y_test, y_pred)
    print(title, "\n", clas)


```

1. Random Forest Classifier


```

In [ ]:
wineRF = RandomForestClassifier()
clas_report(X_train, y_train, X_test, y_test, wineRF, 'Regular Data')
clas_report(X_over_train, y_over_train, X_over_test, y_over_test, wineRF, 'Oversampled Data')
clas_report(X_under_train, y_under_train, X_under_test, y_under_test, wineRF, 'Undersampled Data')
clas_report(X_smote_train, y_smote_train, X_smote_test, y_smote_test, wineRF, 'SMOTE Data')
clas_report(X_impute_train, y_impute_train, X_impute_test, y_impute_test, wineRF, 'Imputed Data')

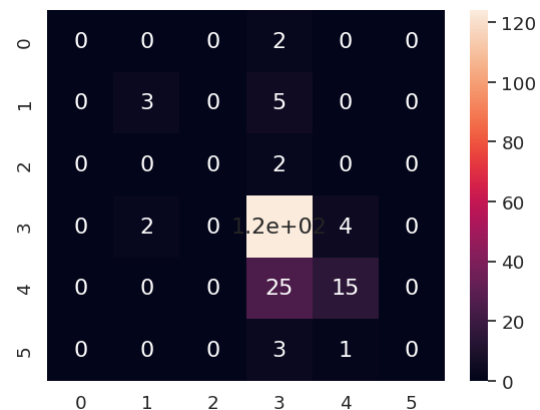
```




Regular Data					
	precision	recall	f1-score	support	
3	0.00	0.00	0.00	2	
4	0.00	0.00	0.00	9	
5	0.74	0.79	0.76	132	
6	0.63	0.72	0.68	127	
7	0.68	0.50	0.58	42	
8	0.00	0.00	0.00	8	
accuracy			0.68	320	
macro avg	0.34	0.34	0.34	320	
weighted avg	0.64	0.68	0.66	320	




Oversampled Data					
	precision	recall	f1-score	support	
3	0.99	1.00	1.00	143	
4	0.00	0.00	0.00	13	
5	0.72	0.75	0.73	131	
6	0.63	0.73	0.68	123	
7	0.82	0.57	0.68	40	
8	1.00	0.25	0.40	4	
accuracy			0.78	454	
macro avg	0.69	0.55	0.58	454	
weighted avg	0.77	0.78	0.77	454	



Undersampled Data					
	precision	recall	f1-score	support	
3	0.00	0.00	0.00	2	
4	0.60	0.38	0.46	8	
5	0.00	0.00	0.00	2	
6	0.77	0.95	0.85	130	
7	0.75	0.38	0.50	40	
8	0.00	0.00	0.00	4	
accuracy			0.76	186	
macro avg	0.35	0.28	0.30	186	
weighted avg	0.73	0.76	0.72	186	



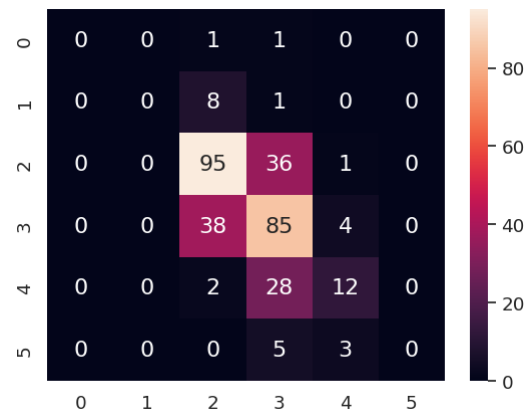
SMOTE Data				
	precision	recall	f1-score	support
3	0.97	1.00	0.98	131
4	0.86	0.93	0.90	123
5	0.72	0.65	0.68	131
6	0.66	0.58	0.61	151
7	0.83	0.90	0.87	147
8	0.94	0.97	0.95	135
accuracy			0.83	818
macro avg	0.83	0.84	0.83	818
weighted avg	0.83	0.83	0.83	818




Imputed Data				
	precision	recall	f1-score	support
0	0.00	0.00	0.00	1
1	0.00	0.00	0.00	1
2	0.00	0.00	0.00	1
3	0.00	0.00	0.00	2
4	0.00	0.00	0.00	9
5	0.41	1.00	0.58	132
6	0.00	0.00	0.00	127
7	0.00	0.00	0.00	42
8	0.00	0.00	0.00	8
9	0.00	0.00	0.00	1
10	0.00	0.00	0.00	1
accuracy			0.41	325
macro avg	0.04	0.09	0.05	325
weighted avg	0.16	0.41	0.23	325

2. Support Vector Machine

```
In [ ]: wineSVM = SVC()
clas_report(X_train, y_train, X_test, y_test, wineSVM, 'Regular Data')
clas_report(X_over_train, y_over_train, X_over_test, y_over_test, wineSVM, 'Oversampled Data')
clas_report(X_under_train, y_under_train, X_under_test, y_under_test, wineSVM, 'Undersampled Data')
clas_report(X_smote_train, y_smote_train, X_smote_test, y_smote_test, wineSVM, 'SMOTE Data')
clas_report(X_impute_train, y_impute_train, X_impute_test, y_impute_test, wineSVM, 'Imputed Data')
```



Regular Data				
	precision	recall	f1-score	support
3	0.00	0.00	0.00	2
4	0.00	0.00	0.00	9
5	0.66	0.72	0.69	132
6	0.54	0.67	0.60	127
7	0.60	0.29	0.39	42
8	0.00	0.00	0.00	8
accuracy			0.60	320
macro avg	0.30	0.28	0.28	320
weighted avg	0.57	0.60	0.57	320




Oversampled Data

	precision	recall	f1-score	support
3	0.92	1.00	0.96	143
4	0.00	0.00	0.00	13
5	0.66	0.73	0.69	131
6	0.57	0.63	0.60	123
7	0.89	0.42	0.58	40
8	0.00	0.00	0.00	4
accuracy			0.73	454
macro avg	0.51	0.46	0.47	454
weighted avg	0.72	0.73	0.72	454




Undersampled Data

	precision	recall	f1-score	support
3	0.00	0.00	0.00	2
4	0.00	0.00	0.00	8
5	0.00	0.00	0.00	2
6	0.73	0.98	0.84	130
7	0.80	0.20	0.32	40
8	0.00	0.00	0.00	4
accuracy			0.73	186
macro avg	0.25	0.20	0.19	186
weighted avg	0.68	0.73	0.65	186



SMOTE Data


	precision	recall	f1-score	support
3	0.89	1.00	0.94	131
4	0.70	0.72	0.71	123
5	0.49	0.45	0.47	131
6	0.48	0.41	0.44	151
7	0.77	0.60	0.67	147
8	0.74	0.99	0.85	135
accuracy			0.69	818
macro avg	0.68	0.70	0.68	818
weighted avg	0.68	0.69	0.68	818




Imputed Data				
	precision	recall	f1-score	support
0	0.00	0.00	0.00	1
1	0.00	0.00	0.00	1
2	0.00	0.00	0.00	1
3	0.00	0.00	0.00	2
4	0.00	0.00	0.00	9
5	0.41	1.00	0.58	132
6	0.00	0.00	0.00	127
7	0.00	0.00	0.00	42
8	0.00	0.00	0.00	8
9	0.00	0.00	0.00	1
10	0.00	0.00	0.00	1
accuracy			0.41	325
macro avg	0.04	0.09	0.05	325
weighted avg	0.16	0.41	0.23	325

3. Artificial Neural Network

```
In [ ]: wineMLP = MLPClassifier()
clas_report(X_train, y_train, X_test, y_test, wineMLP, 'Regular Data')
clas_report(X_over_train, y_over_train, X_over_test, y_over_test, wineMLP, 'Oversampled Data')
clas_report(X_under_train, y_under_train, X_under_test, y_under_test, wineMLP, 'Undersampled Data')
clas_report(X_smote_train, y_smote_train, X_smote_test, y_smote_test, wineMLP, 'SMOTE Data')
clas_report(X_impute_train, y_impute_train, X_impute_test, y_impute_test, wineMLP, 'Imputed Data')
```




Regular Data				
	precision	recall	f1-score	support
3	0.00	0.00	0.00	2
4	0.00	0.00	0.00	9
5	0.66	0.71	0.69	132
6	0.58	0.62	0.60	127
7	0.58	0.50	0.54	42
8	0.00	0.00	0.00	8
accuracy			0.61	320
macro avg	0.30	0.31	0.30	320
weighted avg	0.58	0.61	0.59	320




Oversampled Data				
	precision	recall	f1-score	support
3	0.98	1.00	0.99	143
4	1.00	0.08	0.14	13
5	0.63	0.72	0.67	131
6	0.57	0.59	0.58	123
7	0.62	0.45	0.52	40
8	0.00	0.00	0.00	4
accuracy			0.72	454
macro avg	0.63	0.47	0.48	454
weighted avg	0.73	0.72	0.71	454



Undersampled Data				
	precision	recall	f1-score	support
3	0.00	0.00	0.00	2
4	0.50	0.12	0.20	8
5	0.00	0.00	0.00	2
6	0.78	0.95	0.86	130
7	0.78	0.45	0.57	40
8	0.00	0.00	0.00	4
accuracy			0.77	186
macro avg	0.34	0.25	0.27	186
weighted avg	0.73	0.77	0.73	186




SMOTE Data				
	precision	recall	f1-score	support
3	0.93	0.99	0.96	131
4	0.73	0.79	0.76	123
5	0.56	0.47	0.51	131
6	0.49	0.48	0.48	151
7	0.75	0.67	0.71	147
8	0.85	0.99	0.91	135
accuracy			0.72	818
macro avg	0.72	0.73	0.72	818
weighted avg	0.71	0.72	0.72	818




Imputed Data	precision	recall	f1-score	support
0	0.00	0.00	0.00	1
1	0.00	0.00	0.00	1
2	0.00	0.00	0.00	1
3	0.00	0.00	0.00	2
4	0.00	0.00	0.00	9
5	0.46	0.90	0.61	132
6	0.40	0.15	0.22	127
7	0.30	0.14	0.19	42
8	0.00	0.00	0.00	8
9	0.00	0.00	0.00	1
10	0.00	0.00	0.00	1
accuracy			0.44	325
macro avg	0.11	0.11	0.09	325
weighted avg	0.38	0.44	0.36	325

4. Logistic Regression

```
In [ ]: wineLR = LogisticRegression()
clas_report(X_train, y_train, X_test, y_test, wineLR, 'Regular Data')
clas_report(X_over_train, y_over_train, X_over_test, y_over_test, wineLR, 'Oversampled Data')
clas_report(X_under_train, y_under_train, X_under_test, y_under_test, wineLR, 'Undersampled Data')
clas_report(X_smote_train, y_smote_train, X_smote_test, y_smote_test, wineLR, 'SMOTE Data')
clas_report(X_impute_train, y_impute_train, X_impute_test, y_impute_test, wineLR, 'Imputed Data')
```




Regular Data	precision	recall	f1-score	support
3	0.00	0.00	0.00	2
4	0.00	0.00	0.00	9
5	0.64	0.74	0.69	132
6	0.52	0.58	0.55	127
7	0.52	0.31	0.39	42
8	0.00	0.00	0.00	8
accuracy			0.58	320
macro avg	0.28	0.27	0.27	320
weighted avg	0.54	0.58	0.55	320



Oversampled Data

	precision	recall	f1-score	support
3	0.85	1.00	0.92	143
4	0.00	0.00	0.00	13
5	0.69	0.66	0.67	131
6	0.56	0.62	0.59	123
7	0.64	0.35	0.45	40
8	0.00	0.00	0.00	4
accuracy			0.70	454
macro avg	0.45	0.44	0.44	454
weighted avg	0.67	0.70	0.68	454




Undersampled Data

	precision	recall	f1-score	support
3	0.00	0.00	0.00	2
4	0.50	0.12	0.20	8
5	0.00	0.00	0.00	2
6	0.73	0.92	0.82	130
7	0.45	0.23	0.30	40
8	0.00	0.00	0.00	4
accuracy			0.70	186
macro avg	0.28	0.21	0.22	186
weighted avg	0.63	0.70	0.64	186



SMOTE Data

	precision	recall	f1-score	support
3	0.73	0.89	0.80	131
4	0.52	0.37	0.43	123
5	0.45	0.47	0.46	131
6	0.35	0.30	0.32	151
7	0.43	0.43	0.43	147
8	0.54	0.65	0.59	135
accuracy			0.51	818
macro avg	0.50	0.52	0.51	818
weighted avg	0.50	0.51	0.50	818



Imputed Data	precision	recall	f1-score	support
0	0.00	0.00	0.00	1
1	0.00	0.00	0.00	1
2	0.00	0.00	0.00	1
3	0.00	0.00	0.00	2
4	0.00	0.00	0.00	9
5	0.46	0.77	0.57	132
6	0.36	0.06	0.11	127
7	0.23	0.45	0.31	42
8	0.00	0.00	0.00	8
9	0.00	0.00	0.00	1
10	0.00	0.00	0.00	1
accuracy			0.39	325
macro avg	0.10	0.12	0.09	325
weighted avg	0.36	0.39	0.31	325

4. Hyperparameter Tuning

1. Random Forest Classifier

```
In [ ]: def grid_search(X_train, y_train, model, param_grid, title):
        grid = GridSearchCV(estimator=model, param_grid=param_grid, cv=5)
        grid.fit(X_train, y_train)
        print(title, ":", grid.best_params_)
        print("best score: ", grid.best_score_)

In [ ]: wineRFC = RandomForestClassifier(random_state=417)

param_gridRFC = {
    'n_estimators': [100, 200, 500, 1000],
    'max_depth': [5, 6, 7, 8],
}

grid_search(X_train, y_train, wineRFC, param_gridRFC, 'Regular Data')
grid_search(X_over_train, y_over_train, wineRFC, param_gridRFC, 'Oversampled Data')
grid_search(X_under_train, y_under_train, wineRFC, param_gridRFC, 'Undersampled Data')
grid_search(X_smote_train, y_smote_train, wineRFC, param_gridRFC, 'SMOTE Data')
grid_search(X_impute_train, y_impute_train, wineRFC, param_gridRFC, 'Imputed Data')

Regular Data : {'max_depth': 8, 'n_estimators': 100}
best score: 0.650545343137255
Oversampled Data : {'max_depth': 8, 'n_estimators': 200}
best score: 0.7511049556504104
Undersampled Data : {'max_depth': 7, 'n_estimators': 1000}
best score: 0.7264556502811537
SMOTE Data : {'max_depth': 8, 'n_estimators': 200}
best score: 0.784578819937152
Imputed Data : {'max_depth': 8, 'n_estimators': 100}
best score: 0.6487810068093385
```

2. Support Vector Machine

```
In [ ]: wineSVM = SVC()
param_grid_svm = {
    'kernel': ['rbf', 'sigmoid', 'poly'],
    'C': [1, 10, 100, 1000]
}

grid_search(X_train, y_train, wineSVM, param_grid_svm, 'Regular Data')
grid_search(X_over_train, y_over_train, wineSVM, param_grid_svm, 'Oversampled Data')
grid_search(X_under_train, y_under_train, wineSVM, param_grid_svm, 'Undersampled Data')
grid_search(X_smote_train, y_smote_train, wineSVM, param_grid_svm, 'SMOTE Data')
grid_search(X_impute_train, y_impute_train, wineSVM, param_grid_svm, 'Imputed Data')

Regular Data : {'C': 1, 'kernel': 'rbf'}
best score: 0.6067493872549019
Oversampled Data : {'C': 10, 'kernel': 'rbf'}
best score: 0.7252278025005298
Undersampled Data : {'C': 1, 'kernel': 'rbf'}
best score: 0.7156992563032831
SMOTE Data : {'C': 1000, 'kernel': 'rbf'}
best score: 0.8252764235637917
Imputed Data : {'C': 1, 'kernel': 'rbf'}
best score: 0.6043865515564203
```

3. Artificial Neural Network

```
In [ ]: wineANN = MLPClassifier(activation='logistic')
param_grid_ann = {
    'learning_rate_init': [0.001, 0.01, 0.1],
    'max_iter': [100, 300, 500],
    'hidden_layer_sizes': [(10,), (100,), (300,)]
}
```

```

grid_search(X_train, y_train, wineANN, param_grid_ann, 'Regular Data')
grid_search(X_over_train, y_over_train, wineANN, param_grid_ann, 'Oversampled Data')
grid_search(X_under_train, y_under_train, wineANN, param_grid_ann, 'Undersampled Data')
grid_search(X_smote_train, y_smote_train, wineANN, param_grid_ann, 'SMOTE Data')
grid_search(X_impute_train, y_impute_train, wineANN, param_grid_ann, 'Imputed Data')

Regular Data : {'hidden_layer_sizes': (100,), 'learning_rate_init': 0.01, 'max_iter': 100}
best score: 0.6130024509803922
Oversampled Data : {'hidden_layer_sizes': (300,), 'learning_rate_init': 0.01, 'max_iter': 500}
best score: 0.7411845730027549
Undersampled Data : {'hidden_layer_sizes': (100,), 'learning_rate_init': 0.1, 'max_iter': 100}
best score: 0.7291583529838563
SMOTE Data : {'hidden_layer_sizes': (300,), 'learning_rate_init': 0.01, 'max_iter': 500}
best score: 0.8405772463951371
Imputed Data : {'hidden_layer_sizes': (300,), 'learning_rate_init': 0.1, 'max_iter': 300}
best score: 0.6082715223735409

```

4. Logistic Regression

```

In [ ]: wineLR = LogisticRegression()
param_grid_lr = {
    'penalty': ['l1', 'l2'],
    'solver': ['lbfgs', 'liblinear', 'saga'],
    'max_iter': [100, 300, 500]
}

grid_search(X_train, y_train, wineLR, param_grid_lr, 'Regular Data')
grid_search(X_over_train, y_over_train, wineLR, param_grid_lr, 'Oversampled Data')
grid_search(X_under_train, y_under_train, wineLR, param_grid_lr, 'Undersampled Data')
grid_search(X_smote_train, y_smote_train, wineLR, param_grid_lr, 'SMOTE Data')
grid_search(X_impute_train, y_impute_train, wineLR, param_grid_lr, 'Imputed Data')

Regular Data : {'max_iter': 100, 'penalty': 'l2', 'solver': 'lbfgs'}
best score: 0.6004993872549019
Oversampled Data : {'max_iter': 100, 'penalty': 'l1', 'solver': 'saga'}
best score: 0.6690521599612509
Undersampled Data : {'max_iter': 100, 'penalty': 'l2', 'solver': 'lbfgs'}
best score: 0.7088971521857428
SMOTE Data : {'max_iter': 100, 'penalty': 'l2', 'solver': 'lbfgs'}
best score: 0.5437655422397685
Imputed Data : {'max_iter': 100, 'penalty': 'l2', 'solver': 'lbfgs'}
best score: 0.5981608706225681

```

5. Testing Final Models

```

In [ ]: def final_report(X_train, y_train, X_test, y_test, model, title):
    model.fit(X_train, y_train)
    y_pred = model.predict(X_test)
    print(title, ":", model.score(X_test, y_test))
    # clas = classification_report(y_pred, y_test)
    # print(clas)

```

1. Random Forest Classifier

```

In [ ]: wineRF_regular_final = RandomForestClassifier(n_estimators=100, max_depth=8, random_state=417)
wineRF_over_final = RandomForestClassifier(n_estimators=200, max_depth=8, random_state=417)
wineRF_under_final = RandomForestClassifier(n_estimators=1000, max_depth=7, random_state=417)
wineRF_smote_final = RandomForestClassifier(n_estimators=200, max_depth=8, random_state=417)
wineRF_impute_final = RandomForestClassifier(n_estimators=100, max_depth=8, random_state=417)

final_report(X_train, y_train, X_test, y_test, wineRF_regular_final, 'Regular Data')
final_report(X_over_train, y_over_train, X_over_test, y_over_test, wineRF_over_final, 'Oversampled Data')
final_report(X_under_train, y_under_train, X_under_test, y_under_test, wineRF_under_final, 'Undersampled Data')
final_report(X_smote_train, y_smote_train, X_smote_test, y_smote_test, wineRF_smote_final, 'SMOTE Data')
final_report(X_impute_train, y_impute_train, X_impute_test, y_impute_test, wineRF_impute_final, 'Imputed Data')

Regular Data : 0.6375
Oversampled Data : 0.748898678414097
Undersampled Data : 0.7419354838709677
SMOTE Data : 0.7481662591687042
Imputed Data : 0.40923076923076923

```

2. Support Vector Machine

```

In [ ]: wineSVM_regular_final = SVC(kernel = 'rbf', C = 1)
wineSVM_over_final = SVC(kernel = 'rbf', C = 10)
wineSVM_under_final = SVC(kernel = 'rbf', C = 1)
wineSVM_smote_final = SVC(kernel = 'rbf', C = 1000)
wineSVM_impute_final = SVC(kernel = 'rbf', C = 1)

final_report(X_train, y_train, X_test, y_test, wineSVM_regular_final, 'Regular Data')
final_report(X_over_train, y_over_train, X_over_test, y_over_test, wineSVM_over_final, 'Oversampled Data')
final_report(X_under_train, y_under_train, X_under_test, y_under_test, wineSVM_under_final, 'Undersampled Data')
final_report(X_smote_train, y_smote_train, X_smote_test, y_smote_test, wineSVM_smote_final, 'SMOTE Data')
final_report(X_impute_train, y_impute_train, X_impute_test, y_impute_test, wineSVM_impute_final, 'Imputed Data')

Regular Data : 0.6
Oversampled Data : 0.7268722466960352
Undersampled Data : 0.7311827956989247
SMOTE Data : 0.8092909535452323
Imputed Data : 0.40615384615384614

```

3. Artificial Neural Network

```

In [ ]: wineANN_regular_final = MLPClassifier(learning_rate_init=0.01, max_iter=100, hidden_layer_sizes=(100,))
wineANN_over_final = MLPClassifier(learning_rate_init=0.01, max_iter=500, hidden_layer_sizes=(300,))
wineANN_under_final = MLPClassifier(learning_rate_init=0.1, max_iter=100, hidden_layer_sizes=(100,))
wineANN_smote_final = MLPClassifier(learning_rate_init=0.01, max_iter=500, hidden_layer_sizes=(300,))
wineANN_impute_final = MLPClassifier(learning_rate_init=0.1, max_iter=300, hidden_layer_sizes=(300,))

final_report(X_train, y_train, X_test, y_test, wineANN_regular_final, 'Regular Data')
final_report(X_over_train, y_over_train, X_over_test, y_over_test, wineANN_over_final, 'Oversampled Data')
final_report(X_under_train, y_under_train, X_under_test, y_under_test, wineANN_under_final, 'Undersampled Data')
final_report(X_smote_train, y_smote_train, X_smote_test, y_smote_test, wineANN_smote_final, 'SMOTE Data')
final_report(X_impute_train, y_impute_train, X_impute_test, y_impute_test, wineANN_impute_final, 'Imputed Data')

Regular Data : 0.578125
Oversampled Data : 0.7290748898678414
Undersampled Data : 0.7150537634408602
SMOTE Data : 0.8398533007334963
Imputed Data : 0.12923076923076923

```


4. Logistic Regression

```
In [ ]: wineLR_regular_final = LogisticRegression(penalty='l2', solver='lbfgs', max_iter=100)
wineLR_over_final = LogisticRegression(penalty='l1', solver='saga', max_iter=100)
wineLR_under_final = LogisticRegression(penalty='l2', solver='lbfgs', max_iter=100)
wineLR_smote_final = LogisticRegression(penalty='l2', solver='lbfgs', max_iter=100)
wineLR_impute_final = LogisticRegression(penalty='l2', solver='lbfgs', max_iter=100)

final_report(X_train, y_train, X_test, y_test, wineLR_regular_final, 'Regular Data')
final_report(X_over_train, y_over_train, X_over_test, y_over_test, wineLR_over_final, 'Oversampled Data')
final_report(X_under_train, y_under_train, X_under_test, y_under_test, wineLR_under_final, 'Undersampled Data')
final_report(X_smote_train, y_smote_train, X_smote_test, y_smote_test, wineLR_smote_final, 'SMOTE Data')
final_report(X_impute_train, y_impute_train, X_impute_test, y_impute_test, wineLR_impute_final, 'Imputed Data')

Regular Data : 0.578125
Oversampled Data : 0.7048458149779736
Undersampled Data : 0.6989247311827957
SMOTE Data : 0.5122249388753056
Imputed Data : 0.39384615384615385
```